

Project Report

GitHub URL

https://github.com/ConorSaund/UCDPA_conorsaunders

Abstract

In this project, the analysis was of pricing of Airbnb listings in ten European capital cities. The dataset consisted of twenty files, including weekday and weekend data for each city. The aim was to gain insights into the factors affecting the pricing of Airbnb listings in these cities.

Introduction

The various attributes in the dataset, such as pricing, satisfaction levels, and location to a metro, provide a rich source of information that can be used to extract insights. Working with this dataset requires creative thinking and the application of various analytical techniques to derive meaningful results. We all know the reasons that causes rental spaces to increase their pricing, it is interesting to try to prove it with certain factors.

Dataset

The dataset used was of Airbnb pricing in ten European capital cities. The attributes used in the data sets were room types, cleanliness, distance from the city centre to name a few. The data sets consist of weekday data and weekend data for the ten cities, leading to twenty files to work with. There were approximately 51'700 observations of non-null entries.

The dataset was chosen as it provides an interesting opportunity to analyse the Airbnb market in some of the most popular European tourist destinations. The dataset was also accredited by the National Science centre in Poland under project 2017/27/N/HS4/00951. By examining the pricing trends and characteristics of Airbnb listings in these cities we can gain insights into the broader tourism sector, and with further analysis it could be used to identify future areas of growth.

Implementation Process

The initial code is filled by Kaggle, as I was using their online notebook for coding. It is the default code used to pull datasets straight from Kaggle. The datasets used throughout the assignment were then input using an integration where it took all CSV files available from the domain in Kaggle. Following this, an initialisation of the main packages that were to be used throughout the code was needed.

A load of all of the data that was to be used was then carried out. This consisted of twenty CSV files, of ten European cities of weekday data and weekend. A list was created of each of the links for the CSV files. This was called 'dataframe' for ease of use as it was not going to be used past this point. Before we could concatenate all the files and condense them into one file, we need to ensure that all files have the same number of columns. This was carried out with a 'for' loop to iterate through the files. 'df.shape' was used to output the number of rows and columns in each file. 'df.columns' was also used to output the headers of each file. A quick check was then carried out to ensure the number of columns were the same along with the headers.

A function was later defined to concatenate two CSV files into one file. This was used for the weekday and weekend pricings of each city to be merged into one file for each city and any errors or blanks removed.

The next block of code for "1. Comparison in room types", a comparison was carried out of room types from across all of the cities in the datasets. A grouping was done of each "city" and "room_type", sorting of the room types by count in a descending order, and formatting of the graph. Capitalisation of the city names, along with adding axis labels, and a city by count legend on the right handside of the graph.

A second graph was created to show the mean and median pricing of the house lets. Similar to the formatting of the previous graph, Capitalisation was carried out for the x-axis labels, and a table formed for an easier overview of the pricings.

The second insight used for the datasets of Airbnb lettings amongst ten European cities was Superhost proportions in the lets. The code is to iterate through the cities and find the counts of superhosts. A calculation is made showing the proportion of Superhosts and not. A basic table is printed showing the mean satisfaction level of each "room_type".

A multiple linear regression was carried out on the independent variables of "cleanliness_rating" and "metro_dist" along with a dependent variable of "realSum", to explore the relationship between them. The first step was to check for any missing or error values in the dataset using ".isnull()". This returned a boolean for each cell indication whether it contained a missing value or not. Based on this we went further and removed rows that contained missing cells and did not include them in any calculations. A check for incorrect values was carried out using ".describe()", and further went on to remove any outliers that would skew the data. Calculations of all necessary data were carried out and a regression model used, which produces an output to interpret.

Another linear regression analysis was carried out on the distance from metros, cleanliness rating, number of rooms and the attraction index. The code was to first

calculate the average distance to metros for each room type. A bar plot is then formed for each room type. After the plot was created, a linear regression model was used, OLS – Ordinary Least Squares method. The regression output shows various statistics including the R-squared value which indicates the proportion of the variation in the dependent variable by the independent variable.

The final block was divided into two parts, a heat map and cluster map showing the locations of each Airbnb location as was given in the dataset by longitude and latitude of each. These maps show the distribution of the locations around each city, for which we can determine if locations are in tourist areas, and if this would lead to a higher price on average. The first part imports the relevant packages and creates a centred map of the cities. It creates a MarkerCluster overlay on the map and allows for interactive uses on the map. The second map shows a similar result but leaning on a heatmap overlay.

Results

(Include the charts and describe them)




Figure 1 shows the counts of room types amongst the ten European cities. It is clear that the most common room type was an “Entire home/ apt”, with “ Private room” slightly less and finally “Shared room” with the lowest count.




Figure 2 shows the mean and median prices of the European cities. From the graph and table we can see that Paris has the highest mean price at 392.53, which is higher than most of the other European cities. Athens has the lowest mean price at 151.74, which is significantly lower than the other cities. This suggests that Athens could be a relatively affordable city to visit or rent in. The median price for all cities is greater than the mean signifying a skew towards the higher end. Amsterdam has a significantly higher pricing compared to all of the other cities, showing that accommodation can be quite expensive and driving up the pricing. Lisbon, Barcelona and Vienna all have relatively similar mean and median pricing, which is suggesting the distribution of prices are relatively even in the aforementioned cities. London has a relatively high mean price but lower median. This shows that London has some accommodation which is driving up the average pricing in the area.

```
Proportion of superhosts: False  0.744097
True  0.255903
Name: host_is_superhost, dtype: float64
Mean guest satisfaction by room type:
room_type
Entire home/apt  92.888691
Private room    92.231156
Shared room     89.674863
Name: guest_satisfaction_overall, dtype: float64
(figure 3)
```

Figure 3 shows the proportion of Superhosts amongst the datasets. Working through the output we can see that only approximately 25% of hosts were superhosts. We can see that the majority of the stays were with non-superhosts. We can also see from the output that lets with an "Entire home/ apt" had higher satisfaction ratings than "Private room" or "Shared room".

We do not need further data to understand the reasoning behind the higher satisfaction ratings for "Entire home/ apt" or "Private room". We can use our own personal experiences for this. It is interesting that more customers did not choose lets with Superhosts. Typically a let with a superhost would generally increase the price for the extra care and attention than the customer would receive. It would be an interesting collection of data to see the reasoning behind not selecting a superhost.





Figure 4 shows a bell-shaped histogram which is skewed towards the right. We can see from the counts on the previous figure, Figure 1, that the cheapest room type, “Shared room”, the count was quite low, with “Private room” next and finally “Entire home/ apt” which would skew the graph to the upper end of the pricing. The mode of the graph was approximately 120.00, mean roughly 280.00 and median 220.00.



(figure 5)

Figure 5, show a boxplot of the room types against the pricing. A box plot is a standardised way of displaying the distribution of data based on Q1, Q3, minimum, maximum and median. We can immediately see the number of outliers for all of the room types. All of the boxplots' minimum was roughly the same, with "Shared room"'s maximum being significantly less than the other room types. "Entire home/ apt" maximum was substantially greater. All of the boxplots above are all skewed towards the upper end of the pricing, signifying there are more properties at a higher value than lower.

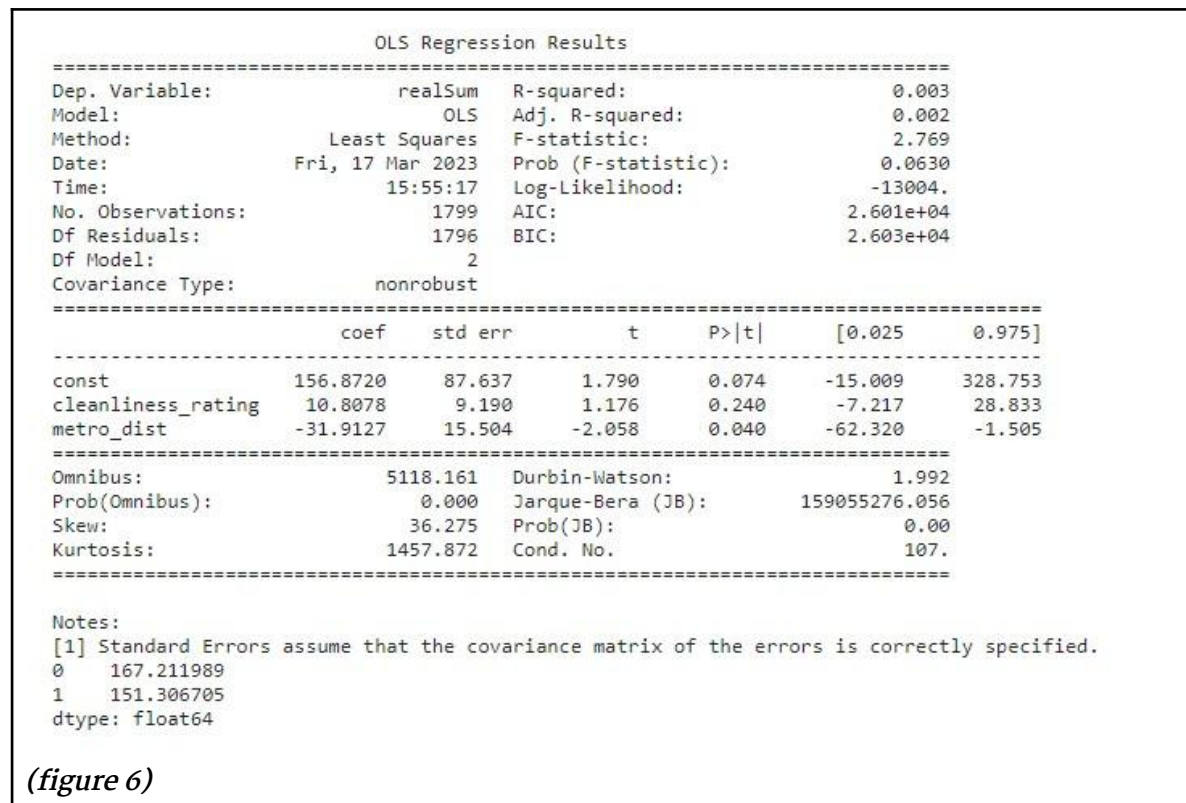


Figure 6 shows the summary output of an OLS (ordinary least squares) linear regression model, which is used to predict rental prices, the dependent variable, based on two independent variables, cleanliness rating and distance to the nearest metro station.

The R-squared value represents the proportion of variance in the rental prices that is explained by the independent variables in the model. In this case, the R-squared is very low at 0.003, indicating that the model is not very effective at explaining the variation in rental prices.

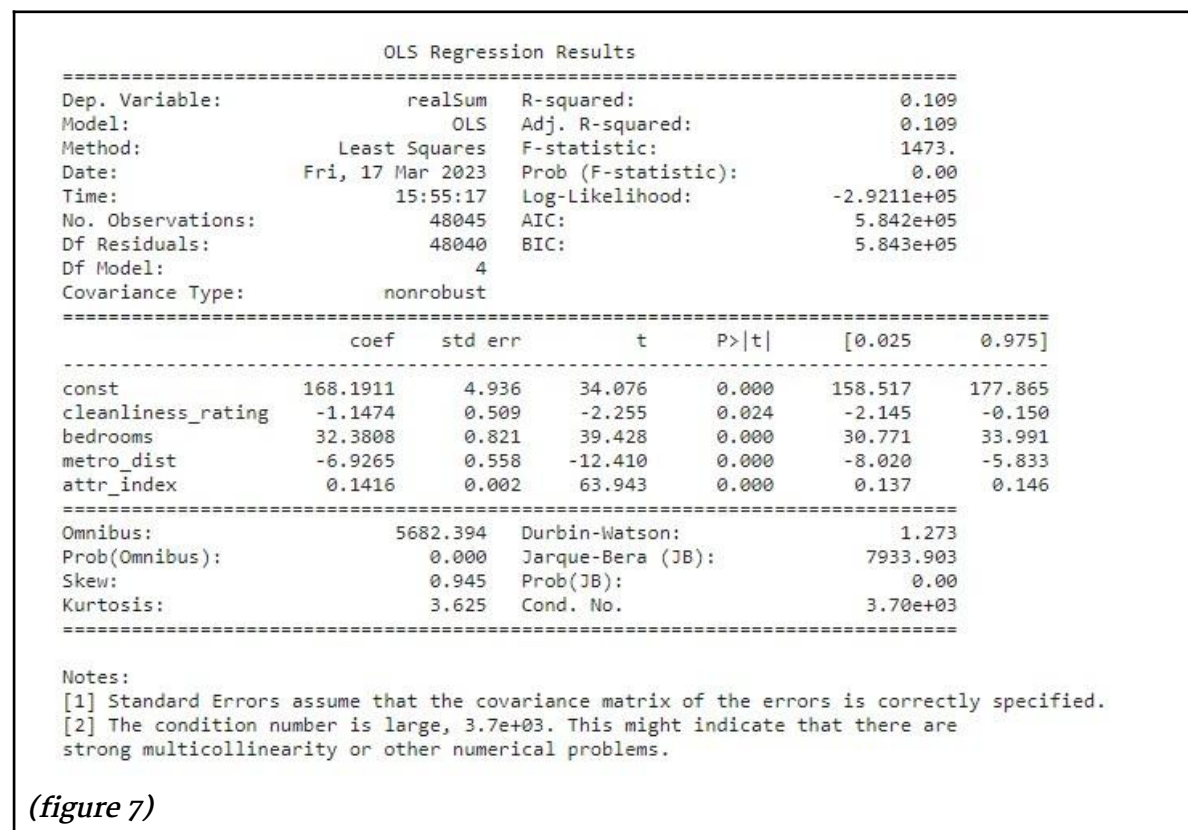
The coefficients represent the estimated effect of each independent variable on rental prices. In this case, the cleanliness rating has a positive coefficient of 10.81, indicating that as cleanliness rating increases, rental prices tend to increase as well. The distance to the nearest metro station has a negative coefficient of -31.91, indicating that as distance to the metro station increases, rental prices tend to decrease.

The P-value indicates the significance of each coefficient, or the likelihood that the observed effect of each independent variable on rental prices is due to chance. In this case, the p-value for cleanliness rating is 0.24, which is not statistically significant at the conventional level of 0.05. The p-value for distance to the metro station is 0.04, which is statistically significant at the 0.05 level.

The confidence interval values indicate the range of values within which the true coefficients are likely to fall with a certain level of confidence. In this case, we can be 95% confident that the true value of the cleanliness rating coefficient falls between -7.22 and

28.83, and that the true value of the distance to the metro station coefficient falls between -62.32 and -1.51.

Finally, the output includes the predicted rental prices for two new observations with cleanliness ratings of 4.5 and 4.8, and metro distances of 1.2 and 1.8. The predicted prices are 167.21 and 151.31, respectively.



The output above is the result of running an OLS (Ordinary Least Squares) regression analysis on the provided data. The model aims to predict the dependent variable "realSum" (i.e., the actual price of the accommodation) based on the independent variables "cleanliness_rating," "bedrooms," "metro_dist," and "attr_index."

The R-squared value of the model is 0.102, which indicates that only 10.2% of the variation in the "realSum" variable is explained by the independent variables included in the model. The adjusted R-squared value is the same, which suggests that adding or removing independent variables from the model does not significantly affect the explanatory power of the model.

The coefficients of the independent variables indicate the direction and magnitude of their effect on the dependent variable. The intercept (or constant) term is 163.9941, which represents the expected value of "realSum" when all independent variables are equal to zero.

The coefficient of "cleanliness_rating" is -0.3427 , which means that a one-unit increase in the cleanliness rating is associated with a decrease of 0.3427 units in "realSum," holding all other variables constant. However, this coefficient is not statistically significant (i.e., $P > |t| > 0.05$), indicating that cleanliness_rating does not have a significant effect on the actual price of the accommodation.

The coefficient of "bedrooms" is 24.4770 , which means that a one-unit increase in the number of bedrooms is associated with an increase of 24.4770 units in "realSum," holding all other variables constant. This coefficient is statistically significant (i.e., $P > |t| < 0.05$), indicating that the number of bedrooms has a significant effect on the actual price of the accommodation.

The coefficient of "metro_dist" is -5.9582 , which means that a one-unit increase in the distance to the nearest metro station is associated with a decrease of 5.9582 units in "realSum," holding all other variables constant. This coefficient is statistically significant (i.e., $P > |t| < 0.05$), indicating that the distance to the nearest metro station has a significant effect on the actual price of the accommodation.

The coefficient of "attr_index" is 0.1269 , which means that a one-unit increase in the attraction index is associated with an increase of 0.1269 units in "realSum," holding all other variables constant. This coefficient is statistically significant (i.e., $P > |t| < 0.05$), indicating that the attraction index has a significant effect on the actual price of the accommodation.

Overall, the results of the regression analysis suggest that the number of bedrooms, the distance to the nearest metro station, and the attraction index are significant predictors of the actual price of the accommodation, while cleanliness rating does not have a significant effect.

Insights

1. By sorting and grouping the room types by city we can see the distribution of the types in each city. This can provide an insight into the most popular types of accommodation in each and how they differ amongst each market.
2. A basic table was output showing the proportion of superhosts in each letting. This can help to identify if there is an influence in using lets from superhosts or if it is an important factor in guest satisfaction and also how it can affect price.
3. By carrying out linear regressions we can explore the relationship between multiple variables and see how they affect the pricing of the lets. The output from the model can be used to make data-driven decisions on price.
4. A plot and regression was used to visualise the relationship between the metro distance and price. This can be useful in identifying which room types are most popular in areas with better access to public transport.

5. The two maps show the distribution of the lettings and can help to identify how location affects pricing and whether it is an important factor for guests when choosing their rental.

References

“The Devastator”, *Airbnb prices in European cities*, Kaggle.

https://zenodo.org/record/4446043#.ZBW_n3bP1D9. Available at:

<https://www.kaggle.com/datasets/thedevastator/airbnb-prices-in-european-cities>

(Accessed: March 2023).